PROBLEM 1 / FIND THE PROMOS (ADS)

PromoScout

An AI-assisted editor tool that turns timestamped microdrama screenplays into source-grounded, ranked 20-60 second promo recommendations.

THE PROBLEM

Editors currently watch every episode and find promo moments by hand. That does not scale as new shows arrive weekly. The output must hook within 3 seconds, end unresolved, and make sense to a new viewer.

WHY THIS PROBLEM

It has a real repetitive workflow, supplied ground-truth timestamps, and a testable output. I scoped the prototype to continuous single-episode clips; stitched promos are a later extension.

01 HOW IT WORKS

1 PARSE	2 GENERATE	3 EVALUATE	4 VALIDATE	5 REVIEW
Read ZIP/TXT, normalize labels, validate every timestamp.	Create 20-60s windows around action, dialogue, reveals, and episode endings.	Gemini scores hook, cliffhanger, standalone clarity, intensity, and risk.	Code rejects invalid cuts, false hooks, weak endings, and duplicates.	Editor inspects source, accepts/rejects, and exports JSON/CSV/Markdown.

02 WHAT AI DOES - AND WHAT IT DOES NOT

PYTHON / RULES	GEMINI	HUMAN EDITOR
Parsing, timestamps, duration, candidate windows, schema validation, penalties, deduplication, diversity, exports.	Semantic judgment: hook, unresolved ending, standalone clarity, emotional/visual potential, warnings.	Checks actual footage, acting, pacing, spoilers, brand safety, and makes the final publishing decision.

03 EXAMPLE OUTPUT

#1 EPISODE 50

00:54.80 - 01:20.00 | 25.2s

Hook: An intimate dialogue and extreme close-up of Catherine and Father Thomas starting the scene with forbidden romantic tension.

Cliffhanger: Agnes discovers the scandalous affair and stumbles back into the shadows in terror right before an abrupt cut to black.

Why it works: The clip features a strong opening hook involving a forbidden relationship, builds immediate romantic and dramatic tension, introduces an unexpected witness, and ends on a high-stakes cliffhanger where a secret is exposed.

Assumptions: screenplay timestamps are authoritative; continuous clips are the MVP; actual footage is unavailable; the system shortlists rather than publishes.

MEASURED ON THE FULL 50-EPISODE PACK

Results, failures, and scale

Every claim below comes from the saved full-pack run and a manual screenplay review of all 15 primary recommendations plus 8 additional examples.

50/50	232	100/100	86.7%
EPISODES PARSED	CANDIDATE WINDOWS	GEMINI OUTPUTS VALID	SHORTLIST-WORTHY TOP 15

04 QUALITY CHECK

MANUAL TOP-15 REVIEW

- 11 strong as-is (73.3%).
- 2 plausible but need boundary/editor review.
- 2 false positives (13.3%).
- 13/15 shortlist-worthy (86.7%).

FULL PIPELINE

- 0 parser errors; 1 retained source warning.
- 65 suitable / 17 borderline / 18 Gemini rejects.
- 35 quality rejections + 8 weaker same-episode duplicates.
- 15 primary promos: exactly 3 from each story arc.

05 WHERE IT FAILS

FALSE CLIFFHANGER	Episode 10 opens with a bloody entrance but ends on the soft question 'Does it hurt?'. Gemini overvalued romantic tension.
FALSE HOOK	Episode 22 ends strongly, but its first 3 seconds are mansion awe and 'This house...' - not a scroll-stopper.
BOUNDARY MISS	A high-scoring Episode 3 fight cut ends before the actual proposal. Gemini correctly rejected it, proving the second stage is necessary.

06 SCALE, COST, AND SUCCESS

SCALE + COST CONTROL

Local rules reduced 232 windows to 100 Gemini calls (top 2 per episode). Responses are cached and resumable. The free tier required 4.5s throttling; the complete run took about 6.5 minutes. Model name and limits are configuration, not code changes.

WHAT I WOULD MEASURE NEXT

Editor acceptance, minutes saved per episode, boundary-adjustment rate, duplicates, and cost per show - then 3-second hold, completion, click-through, and installs versus editor-selected promos in an A/B pilot.

AI USE	Gemini 3.5 Flash Lite provided structured semantic judgments. AI helped build/debug the prototype and draft this document; all metrics come from saved runs. The decisions, manual review, and failure analysis are mine. No claim predicts real views or installs.

Core limitation: screenplay text cannot reveal acting, final footage quality, shot rhythm, audio impact, or platform performance. A real editor remains in the loop.